Courseworkv 1: Experimental Comparison of k-NN and Linear Classification on the Iris data-set

Acereda García, Pablo
19043879

October 18, 2019

This page has been left blank intentionally

Contents

1 Introduction - Machine Learning and Classification Problems

Machine learning, is a field of study which is on the search for predictions towards a certain given input. But the input does not automatically transform into the right answer, by divine intervention. But as a science worthy of its name, it relies on mathematical models which are the result of using statistical models and artificial intelligence (AI).

The process of transformation of a certain input, or training data; to the desired output, is known as classification problem.

In this process, there can be identified three main categories: *supervised learning*, *unsupervised learning* and *reinforcement learning*.

Supervised Learning A problem gets this name when the input goes with its expected output, so the function obtained to classify new inputs comes from labeled training data.

Unsupervised Learning In this case, no data is labeled. On the contrary, unsupervised search helps finding clusters in data (and therefore labelling them) or to find correlation in sets of possibly correlated observations.

Reinforcement Learning The objective of these algorithms is to maximize some notion of cumulative reward, given a certain enviroment and deciding which actions should take a software agent[1].

The scope of this assignment is not to delve into unsupervised learning; nor is to go into detail about reinforcement learning (which has not even been mentioned in the lectures so far).

In the following sections of the project, it is going to be discussed about the dataset used; give a description of the algorithms exploited; and comment how those models are applied to the given dataset.

1.1 Dataset - Iris Dataset

Also known as *Iris flower data set* or *Fisher's Iris data set* (name given after Ronald Fishers, who used the dataset in his publication in 1936); is a dataset formed by three flowers from the same species: **Iris flowers**.

According to the original paper, *"All three species were collected in the Gaspé Peninsula, from the same pasture, picked on the same day and measured at the same time by the same person with the same apparatus"*.

The three species that the original paper talks about are:
- Iris Versicolor
- Iris Setosa
- Iris Virginica

---

[1]"Computer program that acts for a user or other program."

(a) Iris Versicolor   (b) Iris Setosa   (c) Iris Viginica

There were measured the characteristics of 50 flowers from each particular specie; taking into account the length and the width from sepals and petals. The unit of measure used was centimiters (cm).
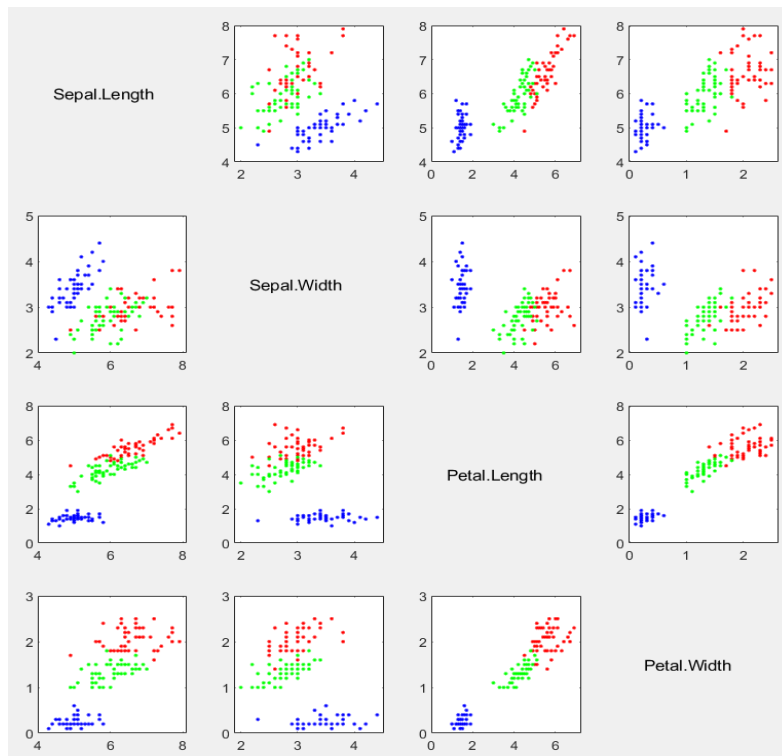
Figure 2: Attributes Relation.

In this specific project, the Iris Dataset has been used to train and test kNN and linear classification, randomly ordering the samples to later apply cross-valitadion method k-fold.

1.2 Models - kNN and Linear Classification

1.2.1 K-Nearest-Neighbors

To classify a new input, by checking which are the closest neighbors; that is the scope this classification method uses. It is possible to understand how it works at a glance, when the model is trained by labeled data, it can predict the class/specie/category of the sample just by computing the k-nearest neighbours; using the *euclidean distance*.

For two dimensions, the formula goes as follows:

$$\sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2} \tag{1}$$

When used in MATLAB, it allows multiple options.

```
Mdl = fitcknn(Tbl, ResponseVarName/formula/Y)
Mdl = fitcknn(X, Y)
Mdl = fitcknn(__, Name, Value)
```

To ease things, *Tbl and X* are attributes from the samples of data used to train the model; *ResponseVarName and Y* are the desired output for the clusters of samples.

formula intends to give an equation on how *Y* depends on the predictor variables.

A very important argument, or list of arguments, which can be specified after any of the afore-mentioned variables is the *Name-Value* Pair Arguments. They allow to edit properties like what breaks ties in case several samples have the same properties; the value for k (how many neighbors are selected to decide the specie), . . .

1.2.2 Linear Classification

There are actually 2 models which create a linear classification: *fitclinear* and *fitcecoc*. As the first of the does not allow multiclass models such the needed for Iris Dataset, that is the one to be described in this section.

In any case, they both use linear equations to "separate" the points from each specie. If a given sample is in the area of a certain specie, the model decides it belongs to that same specie.

$$y = a * bx \tag{2}$$

```
Mdl = fitcecoc(Tbl, ResponseVarName/formula/Y)
Mdl = fitcecoc(X, Y)
Mdl = fitcecoc(__, Name, Value)

[Mdl, HyperparameterOptimizationResults] = fitcecoc(__)
```

At first sight it can be seen that both linear and kNN classification have very similar parameters.

The diference can be found for eample in the pair of arguments *Name-Value*. Can modify the behaviour of the binary learners, design partitions, . . .

There is also one more variable, *HyperparameterOptimizationResults*, which is the description of the cross-validation optimization of hyperparameters. It is not going to be used, as the cross-validation is to be implemented manually.

1.3 Application - Models to Dataset

All these theory is very accurate, and also compelling but, how is this going to help with the Iris Dataset? Fairly simple, by implementin.g a k-fold crossvalidation which will allow to know the accuracy of the models for the given data.

With kNN, the more neighbors that are from the same specie, the more chances that input has to be from that specie. Varying the significant number of neighbors, specified by k might also be helpful when it comes to accuracy.

On the other hand, linear classification uses multiclass linear equations to predict if the given new sample is from one specie or not. In this case, no parameters have been changed.

2 Experimental Work

The aforementioned models have been implemented in MATLAB, in addition to other statistical methods and algorithms in order to predict new samples from the species *Iris Versicolor*, *Iris Setosa* and *Iris Virginica*.

To make it more bearable, the structure of the program could be simplified as follows:

1. K-fold implementation. Henceforth, Data must me splitted in k same-sized pieces.

2. Model generation on training data from *Iris Dataset*, previously separated by k-fold: **fitcknn** and **fitcecoc**.

3. Test models by predicting labels from testing data.

4. Compute error based on the predicted labels and their actual values.

5. Compute the average of all the errors obtained in each iteration performed by the k-fold cross-validation.

6. Data visualization.

2.1 Crossvalidation - K-fold

As it has been mention before, it has been implemented crossvalidation using k-folding. It works by splitting data in k same-sized parts, and iterating through them, so every chunck of data is being used for training and for testing.
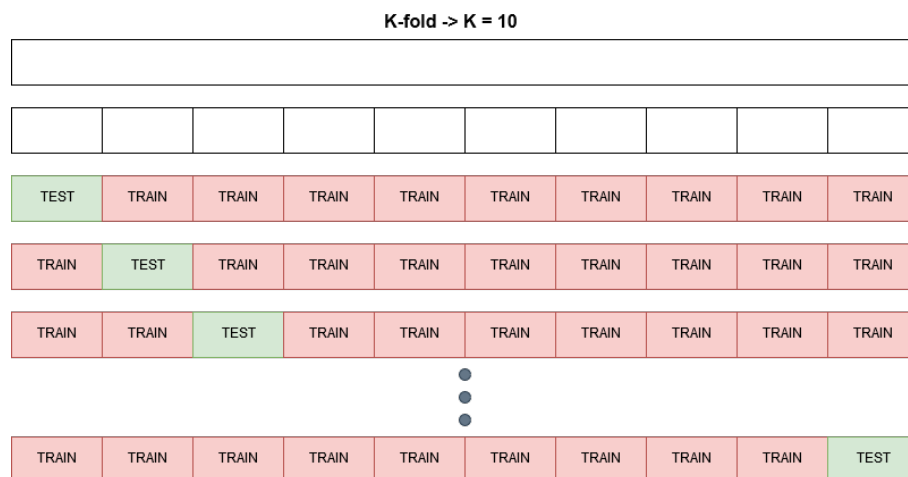

Figure 3: K-fold crossvalidation.

2.2 Modeling and Testing - fitcknn and fitcecoc

The complexity of this procedure is not inherently very high; once data has been separated into chunks from the same size, it is just needed to specify certain parameters at most (1).

It is not needed any forward step to obtain and test the models.

7

2.3 Error - Cost Function

The implementation chosen for this problem is not other than the Mean Error (ME - 3).

$$\frac{1}{N}\sum_{i=1}^{N} I(a,b)$$

$$\begin{cases} I(a,b) = 1, if\, a \neq b \\ I(a,b) = 0, if\, a = b \end{cases} \quad (3)$$

The usage of any other studied error equation would be absurd. For example, taking the Mean Squared Error (MSE - 4):

$$\frac{1}{N}\sum_{i=1}^{N} I(a,b)^2 \tag{4}$$

To square something which has a binary value, 0 or 1^2, would actually result in the same value.

Same happens with the Root Mean Squared Error (RMSE - 5), only in this case it would be with the square root of the same result that would be obtained with ME (Equation 3) but in this case proportional to the square root. Which, at least for the scope of this project is not really necessary.

$$\sqrt{\frac{1}{N}\sum_{i=1}^{N} I(a,b)^2} \tag{5}$$

There is another function already implemented in MATLAB called **resubLoss**, but uses weights in the model and other parameters that are not available, albeit it is more accurate, it cannot be implemented.

3 Conclusion

In brief, it has been used two different classification models. It has also been used k-fold crossvalidation to find the error rate of each model.

Looking at the results, it can be observed that the most suitable value for k-fold in linear classification is surprisingly, 100 (after repeating the program 20 times). The values for k have been changed from 10 to 100 taking steps of 10 at a time. This might make sense; when data is divided into 100 different pieces, the testing chunks are formed by two elements, which rarely gives an error rate higher than 0.

On the other hand, in kNN the value for k in k-fold has been fixed to 10, changing instead the number of neighbors (from 1 to 150). After looking at different executions (once again, 20 runs), the most suitable number of neighbors seems to be 17.

The writer of the article has the beliefs that the selection of the classification method depends on the number of samples in the model mainly: if the number of samples is too large, kNN classification takes too much time to label new elements. On the other hand, if resources and time are not that important, linear classification would mean the loss of classification accuracy for most cases (kNN looses accuracy after the number of neighbors reaches 80, error rate goes through the roof).

---

[2]Should be taken into accout that the predicted label for the sample data can only be right or wrong

Appendix: Code

As any respected MATLAB implementation, first of all is needed to have a clean enviroment:

```matlab
% clear variables
% close opened windows
% clean the comand windows
%------------------------------------------------------------------------%
clear variables;
close all;
clc;
%------------------------------------------------------------------------%
```

Now it is turn to load the data, and order it randomly so that later it is easier to perform the k-fold crossvalidation.

```matlab
% %%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%

% COURSEWORK 1: EXPERIMENTAL COMPARISON OF K-NN AND LINEAR CLASSIFICATION
% ON THE IRIS DATA-SET
% AUTHOR: PABLO ACEREDA

% %%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%

%                              LOAD DATA
%                              ===========
load fisheriris

%                                DATA
%                                ======

% Want to order samples randomly to apply cross-validation
P = randperm(length(species));

% Data from flower (all in cm):
% { Sepal length | Sepal width | Petal length | Petal width }
X = meas(P, :);

% Class label:
% - Setosa
% - Versicolor
% - Virginica
Y = species(P);
```

Once the data is obtained, it might be a good idea to ask the user on the size of the chunks of data. It has also been implemented a log file from which all operations will be recovered post-execution.

```matlab
%                     DATA PARTITIONING INFORMATION
%                     ===============================
```

9

```matlab
kDiv = input(['How many divisions do you want in the sample data ', ...
              '(more that one divisions: (k > 1)?\n']);
N    = length(species);
jump = ceil(N / kDiv);
%                          LOG FILE CREATION
%                          ==================

% File name
filename = 'experiments.log';
% Moves path to file's path
cd(fileparts(mfilename('fullpath')));

% Creates file in case it does not existe
edit(filename);
% Open file
logFile = fopen(filename, 'w');

% Writes in log file
fprintf(logFile,                                    ...
        ['INFO: ',                                  ...
         '==== NEW EXPERIMENT ====\n',              ...
         'INFO: ',                                  ...
         'Data is going to be divided into %i pieces.\n'], ...
        kDiv);
```

   Some other variables and data is needed along the execution of the program. It is a good
practice to initialize the values and the format before they are used.

```matlab
% Valid number of divisions
if kDiv < N

    % Write into log
    fprintf(logFile,                            ...
            ['INFO: ',                          ...
             'The value for k is valid ',       ...
             '(k[%i] < number of samples[%i]) ', ...
             '-> VALID.\n'],                    ...
            kDiv,                               ...
            N);

    %                          VARIABLES
    %                          ===========

    % Later used for the plots
    redRGB   = [255, 0, 0] / 256;
    greenRGB = [0, 255, 0] / 256;
    blueRGB  = [0, 0, 255] / 256;

    colors = [redRGB, greenRGB, blueRGB];
```

```matlab
    % Species available
    speciesStr = ['setosa', 'versicolor', 'virginica'];

    % Each column is a different k value
    k_values = input('Insert values for k in kNN (p.e: [1 3 5]).\n');
    % Write into log
    fprintf(logFile, ...
            ['INFO: ', ...
             'The values used for kNN are: %s.\n'], ...
            mat2str(k_values));

    % Column1: Model
    % Column2: Error
    % Column3: Meas
    % Column4: Species
    lin = {[] [100] [] []};
    knn = cell(length(k_values), 4);
    knn(:, 2) = {[100]};
else
    fprintf(logFile,                                    ...
            ['ERROR: ',                                 ...
             'The value for k is too large ',     ...
             '(k[%i] < number of samples[%i]) ', ...
             '-> VALID.\n'],                            ...
            kDiv,                                       ...
            N);
end
```

Having decided whether the number of divisions for the samples is valid or not, is time to actually split the data. It is done by iterating through the data set. At each iteration, data is separated in training and testing samples.

```matlab
    %                       ========
    %                        K-FOLD
    %                       ========

    for i = 1 : 1 : kDiv

        % Writes in log
        fprintf(logFile,                     ...
                '\n-- K-FOLD: %i/%i --\n\n', ...
                i,                           ...
                kDiv);

        if     i == 1

            testX  = X(      1 : jump, :);
            testY  = Y(      1 : jump, :);
            trainX = X(jump + 1 : end,  :);
```

```
            trainY = Y(jump + 1 : end,   :);

        elseif i * jump <= N && i < kDiv

            testX  = X((i - 1) * jump + 1 : i * jump, :);
            testY  = Y((i - 1) * jump + 1 : i * jump, :);
            trainX = [X(1 : (i - 1) * jump, :) ; X(i * jump + 1: end, :)];
            trainY = [Y(1 : (i - 1) * jump, :) ; Y(i * jump + 1: end, :)];

        else

            testX  = X((i - 1) * jump + 1 : end, :);
            testY  = Y((i - 1) * jump + 1 : end, :);
            trainX = X(1                   : (i - 1) * jump    , :);
            trainY = Y(1                   : (i - 1) * jump    , :);

        end
```

Afterwards, Linear Classification and kNN are estimated with their respective models and errors. In the case of kNN, it is also tested different k values previously selecty by the user. It is also checked whether the model computed is better than the ones already modeled. It is only kept the best model.

After having the best models and their respective data, it is time to graphically represent the results.

```
%                            ==========
%                             PLOTTING
%                            ==========

linX = lin(3);  linY = lin(4);
linX = linX{:}; linY = linY{:};

% Iris Data Graphs
% Rows
for i = 1 : 1 : 4
    % Columns
    for j = 1 : 1 : 8
        % Location in graph
        sp = subplot(4, 8, ((i - 1) * 8 + j));
        % Diagonal
        if mod(i, 4) ~= mod(j, 4)
            % Linear estimations
            if(j <= 4)
                % Process performed to select same color for each label
                for k = 1 : 1 : length(speciesStr)
                    % Which specie should be plotted
                    logic  = strcmp(Y, speciesStr(k));
                    logicX = X(logic == 1, :);
                    % Plot the selected specie
                    scatter(logicX(:, j), logicX(:, i), ...
                            6, colors(k), '.');
```

```matlab
%                        ========
%                         LINEAR
%                        ========

%                          MODEL
Mdl_linear = fitcecoc(trainX, trainY, 'ClassNames', ...
                      {'setosa', 'versicolor', 'virginica'});
%                          TESTING
pred_lin = predict(Mdl_linear, testX);
%                          ERROR
err_lin = costfunction(testY, pred_lin) / length(testY);

% Writes in log
fprintf(logFile,                               ...
        ['// LINEAR \\\\\n',                    ...
         'Model: \n',                          ...
         'Prediction for training: %s\n', ...
         'Error: %f\n'],                       ...
        strjoin(pred_lin),                     ...
        err_lin);

%                 MODEL ERROR COMPARISON
if err_lin < cell2mat(lin(2))

    % Writes in log
    fprintf(logFile,                                   ...
            '--> NEW BEST LINEAR MODEL %f < %f\n', ...
            err_lin,                                   ...
            cell2mat(lin(2)));
    % Saves best model data
    lin = {[Mdl_linear]    [err_lin] ...
           [trainX; testX] [trainY; pred_lin]};
end

% Log: extra separation
fprintf(logFile, '\n');

%                        =====
%                         KNN
%                        =====

% Writes in log
for k = 1 : 1 : length(k_values)

    %                      MODEL
    Mdl_knn = fitcknn (trainX, trainY, ...
                       'NumNeighbors', k_values(k));
    %                      TESTING
    pred_knn = predict(Mdl_knn, testX);
    %                      ERROR
    err_knn = costfunction(testY, pred_knn) / length(testY);
                            13
    % Writes in log
    fprintf(logFile,                       ...
            ['** KNN (k=%i) **\n',       ...
             'Model: \n',                ...
             'Prediction for training: %s\n', ...
```

```matlab
                hold on
                % Repeat process with predicted values
                logic  = strcmp(linY, speciesStr(k));
                logicX = linX(logic == 1, :);
                % Plot estimated labels
                scatter(logicX(:, j), logicX(:, i), ...
                        10, colors(k), 'o');
            end
        else
            for k = 1 : 1 : length(k_values)
                % Data is retrieved
                knnX = knn(k, :);
                knnY = knnX(4); knnY = knnY{:};
                knnX = knnX(3); knnX = knnX{:};

                % Process performed to select same color for each
                % label
                for l = 1 : 1 : length(speciesStr)
                    % Which specie should be plotted
                    logic  = strcmp(Y, speciesStr(l));
                    logicX = X(logic == 1, :);
                    % Plot the selected specie
                    scatter(logicX(:, mod(j, 5) + 1), ...
                            logicX(:, i), ...
                            6, colors(l), '.');
                    hold on
                    % Repeats process with estimated values
                    logic  = strcmp(knnY, speciesStr(l));
                    logicX = knnX(logic == 1, :);
                    % Plot estimated labels
                    scatter(logicX(:, mod(j, 5) + 1), ...
                            logicX(:, i), ...
                            10 + 3 * k, colors(l), 'o');
                end
            end
            hold off
        end
        % Hide legend
        legend('off')
    else
        switch(i)
            case 1
                t = "Sepal.Length";
            case 2
                t = "Sepal.Width";
            case 3
                t = "Petal.Length";
            case 4
                t = "Petal.Width";
```

```
                otherwise
            end
            text(0.1, 0.5, t, "Parent", sp); axis off
        end
    end
```

All models, errors, operations, ... are saved in the log file and printed in the *Command Window* for their visualization.

```
fclose(logFile);
type experiments.log
```

Finally, the cost function implemented to compute the mean error of the predictions made by the models.

```
function cost = costfunction(known, estimated)
    cost = sum(int8(not(strcmp(known, estimated))));
end
```